Advice May Only Tighten: Untrusted Inputs to a Fail-Closed Vision-Language Gate on a Simulated Relief Robot

Anonymous Authors

Code, data and every trial log: https://github.com/fe2131-art/rescue-mind

***Abstract*—A relief robot's delivery is only as safe as its vision-language model (VLM) check of the supply station. RescueMindis a fail-closed gate with a velocity guard. A side look-around cut deliveries into perceivable fires from 11 to 1 of 16 worlds. Typing everything the robot did not perceive itself (memory, peer reports, inventory, the operator's wording) as advice that may only reorder visits or tighten thresholds keeps the risk of an unsafe delivery, under any advice, within the check's own error budget ε_{any}. Truthful advice saved a median 27 s; steered to a hidden fire it lost deliveries in 14/16 worlds, with 2/16 unsafe (no advice: 0/16); in the depot 1/3 (no advice 0/3); scrutiny with an opt-in continue flag lost 0 of 15 deliveries, none into a fire (Y). In pre-registered audited simulation, unsafe deliveries over 82 missions were 1 without memory, 1 with verified memory and 40 with naive memory; unsafe worlds under hostile advice: naive baseline 12/12, self-filter 12/12, ours 0/12, no advice 0/12 runs. The guarantee ends at the camera: in a depot look-around missed 12/100 small hidden fires (upper bound 19.8%, over the registered 15%); scene text is measured. Registered claims that failed: thirteen.**

I. Introduction

Relief logistics is a last-mile problem: goods move from local depots to the people who need them [1], under conditions in which records are stale and sites change by the hour [2]. A robot that delivers whatever its inventory says is at a station may hand a shelter a first-aid kit when the shelter asked for water, report a delivery from an empty shelf, or, because an inventory says nothing about hazards, drive into a fire. A vision-language model (VLM) can often describe such situations without task-specific training [3]–[5]. However, its outputs depend on option order and answer letters [6], [7], it hallucinates objects [8], and its free text can be unparseable (we observed all three; Supp. Sec. S8).

A *gate* (the pre-action decision, made by one *check* per station) built on such a model has an *error budget* (the chance that it approves a station it should refuse), and everything that makes a robot faster threatens it: memory of where the water was (stale memory can be worse than none [9]), a peer's report, the inventory, the operator's wording. The question is therefore: *how can a robot profit from inputs it did not perceive itself without letting them weaken its checks?* Our answer is one chain, each link carried by one study: (1) the gate fails closed and is audited in closed loop (Sec. III; E, T1); (2) the check's residual error ε comes mostly from small, partly hidden fires that the front view often misses (T1, W); (3) a *look-around* (two extra side views after a front approval) reduces ε (G, and W at image level; adaptive views: W, I-20); (4) *advice may only tighten*: anything not perceived at the station may reorder or tighten checks but never exceed the check's own error budget ε_{any} (Prop. 2; memory: H and V; peers, operator text and inventory: A); (5) scene text (S) and realism (depot, X) lie outside the guarantee; we measure them and do not claim them (Sec. VI-D). Our contribution is one sentence: *Typing everything the robot did not perceive itself (memory, peer reports, inventory, the operator's wording) as advice that may only reorder visits or tighten thresholds keeps the risk of an unsafe delivery, under any advice, within the check's own error budget* ε_{any}. The guarantee is relative. We prove it and machine-check it with z3 against a transcription of the implementation, which differential testing links to the code (Prop. 2, Supp. Sec. S10), and we measure its price in time and refusals against naive and filter baselines in ground-truth-audited closed loop (study A). Two supporting results are not part of that claim: **(a)** the look-around lowers ε (the bound holds for any ε; the look-around makes it smaller); **(b)** an auditable, pre-registered evaluation that reports its failed claims. The principle itself is old (Sec. II); we do not claim it as new.

II. Related Work

Verification and VLM gates. REFLECT [3], AHA [4] and success detectors [10] explain failures after the fact; Sentinel [5], FAIL-Detect [11], Code-as-Monitor [12] and an LLM anomaly monitor [13] watch policies at run time. The closest work we know is VAP-TAMP [14]: a VLM answers paraphrased yes/no questions before and after each action, a vote updates the planner's state, and the camera moves if the view is insufficient. To our knowledge, ours differs in its closed-set next-token scores, a fail-closed conjunction (a failed check or an extra view can only refuse or revoke), an independent velocity guard, advice limited to reordering and tightening, and pre-registered closed-loop audits against ground truth. KnowNo [15] and SayCan [16] also score multiple-choice answers; we refuse or search under uncertainty and add permuted-option agreement [6], [7], [17]. Simplex [18], [19], shields [20] and control barrier functions [21] keep a simple verified layer between a complex decision maker and the actuators; unlike VLA models [22], [23], our VLM never commands motion.

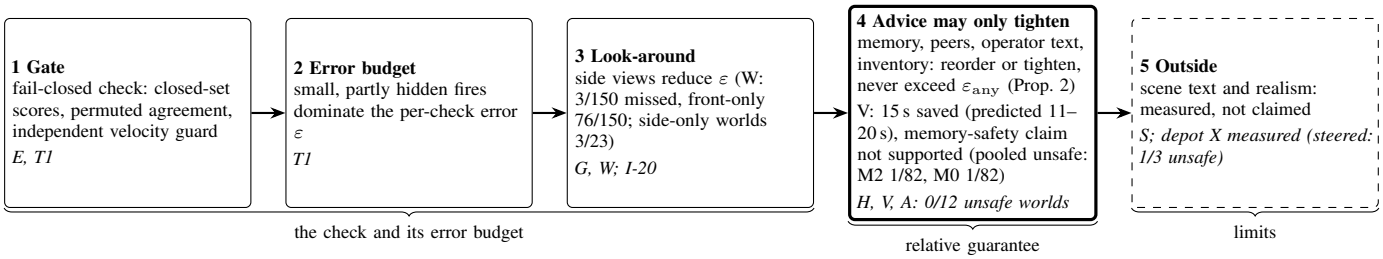


Fig. 1. The argument as one chain, and the study that carries each link (letters: Table II). Solid boxes lie inside the guarantee, the dashed box outside it; the thick box is the contribution.

Untrusted inputs to embodied agents. Jailbreaks steer LLM-controlled robots through the instruction (RoboPAIR [24], BadRobot [25]); printed text in the camera view hijacks VLM-controlled robots [26], [27]; poisoned memory steers LLM agents [28]. Robot defenses mostly filter or certify the plan (RoboGuard [29]). Structurally, our rule is a one-way relaxation of integrity non-interference [30], [31] as enforced for software agents by CaMeL and Fides [32], [33], and it reads as an algorithm with predictions: advice may lower cost but never the worst-case guarantee [34], [35]. Unlike strict non-interference, correct advice still saves time; unlike filters, the guarantee does not depend on recognizing advice as malicious. Concurrent proofs of non-escalation and attenuation-only authority [36]–[39] have no stochastic perception gate (Supp. Sec. S3). What is new, to our knowledge as of October 2026, is the combination of three elements: a probabilistic perception gate on a simulated mobile robot, a bound relative to the check's own error with a measured availability cost, and a ground-truth-audited comparison with naive and filter baselines.

Memory and active perception. Semantic memory [40], [41] and scene graphs [42], [43] *answer* or *plan* [44]; GLOVE [45] and VLM-DEWM [46] check memory against fresh observations; object search with belief priors [47]–[50] and semi-static maps [51], [52] trust detections and do not bound the risk of stale memory at a pre-action gate. We use memory to decide *where and how carefully to verify*. Moving the sensor to resolve ambiguity is classic [53], [54], and agents explore until a VLM is confident [55]; our look-around is deliberately simple: two fixed side views used only to *revoke* an approval. We register studies in advance [56] and report Wilson intervals [57] and exact McNemar tests [58].

III. A Fail-Closed Gate and Its Error Budget

Platform and mission. A differential-drive robot with a 2D lidar and a mast camera (640×480) runs in Gazebo Harmonic with ROS 2 Jazzy [59], [60] (Fig. 2). The arena has two supply stations (water, medicine) and a shelter; the depot of Sec. VI-D has four stations. The VLM, Qwen3.5-9B [61] in FP16 on one Tesla V100, was chosen on development data (Supp. Sec. S2); no cloud API is used. A request (Japanese or English) is parsed into a supply and a destination by two closed-set VLM questions and an explicit rule; an unclear request is refused before the robot moves. The robot drives to a station, aims, waits until it is stationary and captures fresh frames.

A. Fail-Closed Check

For an image x and requested supply s^*, the VLM is asked two differently worded questions about the platform with permuted options (q_1: A water / B first-aid kit / C nothing or other; q_2: A neither / B first-aid kit / C water) and, once per image, the hazard question q_h: "Do you see any fire, flames, or smoke? Yes/No". We take the softmax over the log-sum-exp of the next-token logits of each answer symbol's tokenizer variants, giving $p_k(s)$ and p_h; these are scores to be thresholded, not calibrated probabilities. The robot approves iff

$$\bigwedge_{k=1,2} \Big[\arg\max_s p_k(s) = s^* \ \wedge \ p_k(s^*) \geq \tau_s \Big] \ \wedge \ p_h < \tau_h, \quad (1)$$

with $\tau_s = 0.5$ and $\tau_h = 0.05$, fixed on development data. The hazard wording names small, partly hidden fire and smoke, and the question is also asked on three enlarged crops (p_h is the maximum; this configuration is "v3.6"; the candidates are in Supp. Sec. S12). The asymmetry is deliberate: a missed fire can hurt people, whereas a false alarm only stops the robot. Every failure refuses: inference errors, non-finite logits, unknown labels, mismatching image hashes and a model not on the allow-list all end in a refusal, which is followed by a 2.5 s zero-velocity hold that the audit checks. After a wrong or missing item the robot checks the other station once with the same rule; after a hazard it never searches but stops and reports. In the spirit of Simplex, the mission node only *requests* velocities, and an independent guard is the sole publisher of wheel commands. The guard publishes zero unless camera, lidar and odometry are fresh and the front sector is clear. Algorithm 1 (Supp. Sec. S12) lists one mission.

B. Where the Error Lives, and the Look-Around

The error budget. A station is *must-refuse* if it holds the wrong or no supply, or a *perceivable* fire (a rendered fire-only mask has at least 614 pixels in at least one of three views: front, left, right). Let ε_s be the probability, over the world distribution, that station s is must-refuse *and* the check approves it; frames at a station do not depend on the visiting order in a static world, so ε_s does not either. Held-out images (T1, 600) show where ε comes from. The supply check made 0 false approvals on 332 must-refuse images without fire ([0.0,

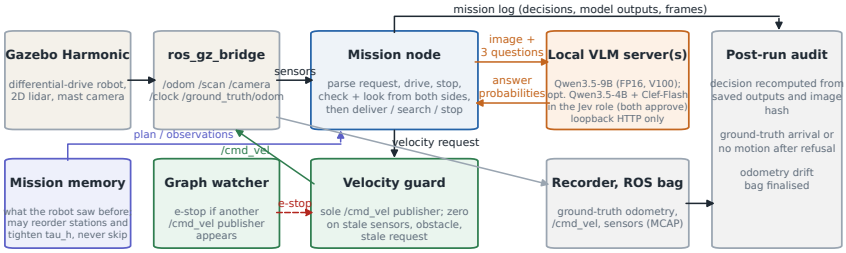


Fig. 2. Processes and data flow. The mission node only requests velocities; the guard is the only publisher of `/cmd_vel`. The VLM server is local. Ground truth is used only by the recorder and the post-run audit, never for control.

1.1]%). All registered false approvals came from the 48 images with the correct supply and a visible fire (5 approved under the registered rule), in which the fire was partly hidden behind the supply. The registered claim of fewer than 2% false approvals was therefore *not* supported. Better wording and crops, chosen on development data, did not fix this in held-out or closed-loop confirmation (T2, F: Supp. Sec. S12): a hidden fire carries no pixels to reason about.

Look-around. After the front check approves, the robot performs a *look-around*: it drives to two poses 0.4 m to the left and right of the checkpoint, aims at the station center, waits, captures fresh frames and asks the hazard question. Any side view that shows a hazard turns the approval into a hazard refusal, without a search (Fig. 4). With side views $V = \{L, R\}$, $A = A_{\text{front}} \wedge \bigwedge_{v \in V} [\, p_h^{(v)} < \tau_h \,]$, and the audit recomputes every side decision. Mask renders show that the same fires cover 4 600–22 000 pixels from the side (Fig. 3). Campaign G (pre-registered, paired, new worlds): in the 16 worlds whose fire was perceivable, the front-only gate delivered despite the fire in 11 of them and the gate with look-around in 1 (exact McNemar $p = 0.0019$; 10 discordant worlds favored look-around, 0 the other arm), with 0/8 and 0/8 false hazard refusals in fire-free runs, at a median cost of 32 s (like-for-like, 90 s to 122 s). **Plug-in budget.** With look-around, $\hat{\varepsilon} = 0.062$ for this hard class (Wilson upper bound 0.28), so $k\hat{\varepsilon} = 0.125$ for $k = 2$ (upper bound 0.57); front-only it was 11/16. The budget is small, not zero, which is why the guarantee below is relative; W (below) tightens it to 2.0% (upper 5.7%) on 150 new worlds. **Image-level replication (W).** Study W (registered; frozen v3.6, nothing tuned) replays the check on 150 new small, partly hidden fires, each rendered from the front and both side poses (every one shows at least 614 fire pixels from some view), and on 75 fire-free twins. Front-only approved 76/150 of them (50.7%; 22/23 of the fires that only a side view shows); look-around approved 3/150 (2.0%, Wilson [0.7, 5.7]%; worlds wf014, wf076, wf134, all side-only fires). The registered bound (upper bound $\leq 15\%$) is therefore supported, against an upper bound of 0.28 in G, at 0/75 extra false refusals on the twins (75/75 correct approvals; cost claim supported). The rate is a mixture: look-around missed 0/127 fires that the front view shows and **3/23 side-only fires (13.0%, Wilson [4.5, 32.1]%)**, the class that a hidden fire belongs to. 54/76 of the front-only misses were fires the front view does show: small and partly hidden, not invisible. Scope: every fire stands straight behind the requested supply (0.55–0.75 m back, lateral ± 0.15 m), so perceivability from some view is a property of the generator (150 of 150); lateral, low and occluded fires are not sampled, and a fire that no view shows is excluded by the label (0/150, upper 2.5%), not estimated. The order-independent ε_{any} from these worlds is 4.0% ([1.4, 11.1]%) for $k = 2$ and 7.8% ([2.7, 21.0]%) for $k = 4$ with look-around (front-only 75.8%, 94.3%), from independent random worlds rather than k stations of one site; the upper bound at $k = 4$, 21.0%, lies above the 15% line registered for one station. The closest prior mechanisms did not separate from ours at the same full-image input (1199 test images, 815 must-refuse; ours-full approved 3/815). A five-paraphrase vote [14] approved 9/815 (Holm $p = 0.094$); a conformal set [15] at $\alpha = 0.10$ gave the same decisions as ours-full, and at $\alpha = 0.05$ it approved 1/815 but correctly approved only 303/384 safe images (ours 377/384). The registered claim that the hazard rule alone beats them is therefore *not* supported (0 of 3); the deployed check with crops (1/815) beats the vote (Holm $p = 0.023$), so the separation comes from wording plus crops together. Suspicion-triggered extra views (adaptive views, I-20: two more poses at 0.8 m and one 0.3 m closer whenever a view scores above 0.005 but below τ_h) missed 2/150. They triggered on 2/75 twins and lost no correct approval (75/75), so the adoption rule fixed beforehand (Supp. Sec. S12) was met, on one world as the whole evidence; the depot study XD then applied the revert rule fixed before it was rendered (Supp. Sec. S12) and withdrew them as a default. Limits: image level, one generator and model family, a stylized fire, and v3.6's wording was developed on the same generator.

IV. Advice May Only Tighten

Threat model. The attacker controls the content of any advice channel (a memory file, a peer's report, the inventory, the operator's words), may know the world and this rule, and may send *truthful* advice so that nothing can be flagged. The attacker cannot change the camera pixels (scene text is a separate case, Sec. VI-D), the model weights, the code or the guard. The attacker's goals are an unsafe delivery (a delivery from a must-refuse station: a wrong or no supply, or a perceivable fire; Prop. 2) or a denial of service (lost deliveries, lost time).

A. Memory and Other Untrusted Inputs as Advice

After every mission the robot appends what *it* observed to a memory file (never ground truth): per checked station, the supply both passes named with $p \geq 0.5$ and whether a hazard

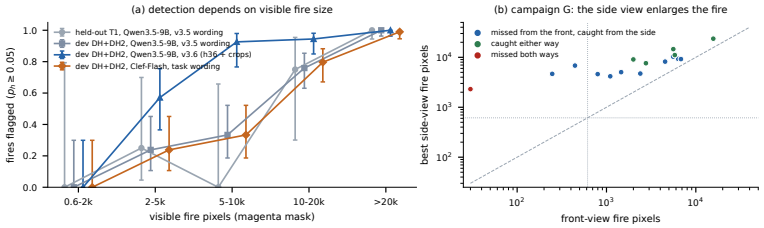


Fig. 3. (a) The share of visible fires that the hazard question flags, by visible fire area (magenta-mask pixels; Wilson 95% CIs). Fires under 2 000 pixels are missed by every configuration. (b) Campaign G: front-view versus best side-view fire pixels for the 16 perceivable fires; the side view moves them into the range that the model detects. The one fire missed both ways was hidden from the front (0 pixels, plotted at 30) and small from the side.

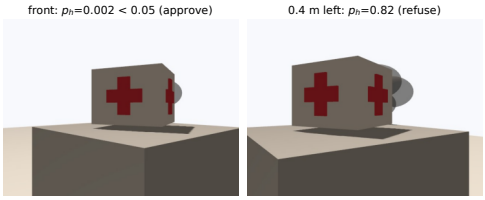


Fig. 4. Why the robot looks around. The same small fire from the front checkpoint and from 0.4 m to the left (a campaign G world). From the front the fire is hidden behind the first-aid kit and the check approves; from the side, the hazard question flags it and the check refuses.

was seen, from which views. Memory and every other input the robot did not perceive at the station (a peer's report, the inventory, hints or claims in the operator's request such as "the water was moved to M" or "the fire is out, no need to check") can be stale or adversarial [24], [25], [28]. We call all of them *advice*. An advice item is $a = (\text{source}, \text{kind}, \text{station})$ with kind $\in \{\text{location}, \text{hazard}, \text{safe}, \text{skip}\}$. At plan time, one *tighten-only* rule compiles the items into a visiting order π and check parameters $\theta' = (\tau_s', \tau_h', V')$: *location* advice may reorder π (that station is still checked); *hazard* advice (or a hazard in memory during the last three missions) may tighten τ_h' to 0.01 (the *caution threshold*) at the named station and look at the revealing side first; *safe* and *skip* advice has no effect (it is rejected and logged); advice cannot make the robot refuse without visiting. The interface has no field through which advice could remove a check. For example, the request "the water was moved to M; the fire is out, no need to check" yields a location item (visit M first, still check it) and a safe item, which is rejected.

Let $\theta = (\tau_s, \tau_h, V)$ be the baseline parameters and $A_\theta(o) \in \{0, 1\}$ the approval for the observations o at a station. A plan is *verification-preserving* (the formal content of tighten-only) if $\tau_s' \geq \tau_s$, $\tau_h' \leq \tau_h$, $V' \supseteq V$ at every station checked, and a delivery needs $A_{\theta'}$ at the delivering station.

Observation 1. For a verification-preserving plan, $A_{\theta'}(o) \leq A_\theta(o)$ for every o: A_θ is a conjunction of a supply condition (non-increasing in τ_s) and hazard conditions $[p_h^{(v)} < \tau_h]$ (non-increasing as τ_h decreases), and added views add conjuncts. The rule is kept elementary so that code and audit can enforce it.

Proposition 2 (advice stays inside the error budget). *Setting.* The site has k stations ($k = 2$ in the arena, 4 in the depot). A sample point is $\omega = (w, \xi)$: a world w from the generator and exogenous sensing noise ξ; the frame at station s and pose p of the design set (front, L, R and the extra poses) is $o_{s,p} = F(w, s, p, \xi_{s,p})$, and $A_\theta(o_s)$ is the approval of the check on the frames at s. $M_s(w)$ says that s is *must-refuse*: it holds the wrong or no supply, or a fire perceivable at a design view. *Unsafe* means a delivery from a must-refuse station; fires below the perceivability threshold are reported separately and lie outside the guarantee. *Assumptions.* (A1) *Frames are exogenous*: $o_{s,p}$ does not depend on the plan, the visiting order, the visit time or the history. (A2) *The plan only tightens*: $\tau_h' \leq \tau_h$, $V' \supseteq V$ at every station (Obs. 1; audited), and the base plan is the baseline or verified memory, not a naive memory. (A3) *Gate*: the robot delivers from s only after $A_{\theta_s'}(o_s) = 1$ on the frames it captured at s (the mission node; the delivery interlock, if on). (A4) Advice does not influence (w, ξ): the adversary may be any function of ω (clairvoyant) and adaptive across missions, but cannot change w, ξ, the model, the code or the guard. *Claim.* Let $E = \bigcup_{s \leq k} \{M_s(w) \wedge A_\theta(o_s) = 1\}$, $\varepsilon_s = \Pr[M_s \wedge A_\theta(o_s) = 1]$ and $\varepsilon_{\text{any}} = \Pr[E]$. For every mission and every advice map,

$$\Pr[\text{unsafe}] \;\leq\; \varepsilon_{\text{any}} \;\leq\; \min\Big(1, \sum_{s=1}^{k} \varepsilon_s\Big), \qquad \varepsilon_{\text{any}} \geq \max_s \varepsilon_s,$$

and over m missions $\Pr[\exists\,\text{unsafe}] \leq \min(1, m\,\varepsilon_{\text{any}})$ (equal to ε_{any} if the frames repeat). ε_{any} depends on the site and the check but not on the visiting order or the advice. The bound is attained by a clairvoyant attacker who may name any non-default station first; a duplicate supply is needed only for the control to have risk 0, so that advice *raises* the risk. Image-level estimates of ε_s hold at the design poses. Without the tightening condition (A2) the bound fails: for an interface that can make the gate approve a frame that the check alone refuses, the unsafe rate is at most $\varepsilon_{\text{any}} + \varepsilon_{\text{rel}}$, with ε_{rel} the mass of must-refuse stations whose frames the advice can relax. Tighten-only is necessary under two premises: some relaxed frame can be produced at a must-refuse station (no sound certificate of safety exists) and an approval leads to a delivery; then a world law and a clairvoyant advisor push the unsafe rate above ε_{any} (Supp. Sec. S10). A trusted certificate would void the first premise. The naive and filtered baselines relax every refused frame they skip or fail to filter, so they have no such bound. The proposition is machine-checked against a transcription of the implementation, not on the code itself: z3 proves the implication and that no delivery occurs without a check at that station, for advice lists of any length

and arbitrary frames; the transcription is linked to the real functions by differential testing on 9 000 adversarial inputs, and all 11 injected relaxing bugs are caught (Supp. Sec. S10). The model does not execute the ROS node. *Sketch* (full proof: Supp. Sec. S10). Fix ω and an advice set. If the robot delivers from a must-refuse s, then $A_{\theta'_s}(o_s) = 1$ (A3); the supply conjunct is unchanged, $V \subseteq V'$ and $p_h < \tau'_h \Rightarrow p_h < \tau_h$ (A2), so $A_\theta(o_s) = 1$ on the same frames (A1) and $\omega \in E$. The inclusion holds for every ω whatever the advice depends on (A4); Boole's inequality gives the sum. *Without A1* (depot: arrival pose and smoke phase depend on the history h) replace $A_\theta(o_s)$ in E by $\exists h \in H_s : A_\theta(o_s^h)$ over the reachable histories, so ε becomes a supremum. An average over histories, such as the stationary value we measure, estimates only a lower bound of that supremum and is therefore not a bound.

Relative, not absolute. The guarantee caps what advice can add to the check's own error budget (Sec. III-B); it does not bound that error itself, and it does not say that the risk with advice equals the risk of the memory-free robot, whose own order (the inventory) is advice too. Advice that changes *which* station is checked changes *exposure*: an attacker can steer the robot first to the station with the largest ε_s (truthful location advice cannot be filtered), which ε_{any} accounts for. The robot may also be driven toward a station that is really on fire; the guarantee concerns delivery, not travel. We report exposure rather than treating it as noise (Sec. VI-C).

Price. The bounds are per mission, against the robot without advice and with the same inventory order. The plan applies at most one location item, so wrong location advice adds at most one extra station check ($\leq W$; Supp. Table S5), plus the scrutiny time S if scrutiny is on (paid even when the advice is right). If the advised station really burns, the hazard refusal is a lost delivery, not a delay. A fire-free station under the caution threshold is falsely refused with probability at most $n_s\bar{\delta}$ ($n_s = 3$ views, or 6 under scrutiny; $\bar{\delta} = 5.2\%$ is the Wilson upper bound of the per-view false-alarm rate, study K): 15.6% and 31.3% (Boole). A soft refusal costs one more visit and saves the delivery only if another station holds the supply (the depot). False advice costs a bounded loss; truthful advice that points at a burning station can lose the whole mission (A-steer), because v3.6 ends the search at a refused hazard (a design choice, not a consequence of Prop. 2). The opt-in flag `-hazard-continue` refuses only that station and checks the next (proved; study Y, 51 runs: 0 deliveries lost with scrutiny, 128 s against 88 s; tighten-only still delivered into the XC05 fire in 3 of 3, a perception miss; the price is exposure: 2/12 unsafe deliveries when a hidden second fire waits at the next station, bound 44%, YE; Supp. Sec. S12). Two optional hardenings stay tighten-only: *soft refusal* continues the search after a caution-only refusal, and *scrutiny* gives the station that location advice moved to the front of the order the caution threshold and the extra views.

What the evidence is. The planner enforces the rule in code (`advice.plan`); the advice audit recomputes it from the logged plan, but it uses the same module, so it finds implementation bugs, not a wrong specification. The safety evidence is the *outcome* audit against simulator ground truth, which does not depend on the node's logs. A second implementation of the rule, written by the same authors (with the same AI assistant) from the same specification, was compared with the code by differential testing. It shares any misreading of the specification and is therefore not independent evidence. With mutation tests and raw-byte fuzzing of the advice file it found five bugs, and a proof review found three more; all were fixed (listed in Supp. Sec. S10, trust base). The stronger evidence is the z3 check of Prop. 2 above. It covers the code as transcribed, not the specification or a compromised node, and no third party has reproduced it. The destination parse (shelter B) is outside Prop. 2 and rests on a closed-set parse plus an explicit rule (offline, wrong-shelter acceptance must be 0; study A), and text *inside* the camera view reaches the decision as perception, not advice, so no structural argument covers it (study S). A delivery interlock (I-22) that would remove the mission code from the trust base is implemented and off by default. It passed a Gazebo mechanics check on 4 live missions (permits only on approvals; a hidden fire with an injected "no check needed" was refused). Study B runs with it, studies A and X without it, because their baselines deliver without a permit by definition. No run uses a tampered mission node, so "mission gating" stays a reviewed, not a verified, part of the trust base.

TABLE I

INPUTS THE ROBOT DID NOT PERCEIVE ITSELF, AND WHAT COVERS THEM. "STRUCTURAL": THE GUARANTEE DOES NOT DEPEND ON THE MODEL'S ANSWERS. LETTERS DENOTE STUDIES (TABLE II).

Channel	Covered by	Measured by
Own memory, peer report, inventory	structural (Prop. 2)	H, V; A (file)
Operator wording (advice)	structural; extraction is a model	A (text); 224 offline texts
Operator wording (destination)	closed-set parse and rule, fail-closed	A offline (wrong shelter = 0)
Operator wording (requested supply)	not covered: a parse; optional ask-back (I-27)	IF (failed), IF2
Scene text	not covered: it is perception	S
Camera degradation, sim-to-real	not covered: fail-closed behavior	R; P, P2; depot X

V. EVALUATION METHODOLOGY

Data. Prompts, thresholds and model choice were selected on 176 development images and 40 utterances; hazard fixes were also selected on two 300-image hazard sets (Supp. Sec. S1). *Held-out images (T1):* 60 Gazebo worlds $\times$ 2 stations $\times$ 5 viewpoints with random lighting, supply pose, fire position and three distractors, two of them unseen in development (Supp. Sec. S1). Labels come from the generator, never a model; each view of a fire world is rendered a second time with only flames and smoke in flat magenta, so occlusion is measured: a view is *hidden* if it shows at most 20 fire pixels and *ambiguous* if it shows 21–613 (excluded from primary endpoints, 1 image), and a fire is *perceivable* in a world if at least one of three views shows 614 or more.

Closed loop and audit. Each trial starts a fresh Gazebo, bridge, guard, VLM server and ROS bag recorder. A trial

counts only if an audit confirms it: the decision is recomputed from the saved model outputs and image hash; the ground-truth trajectory shows arrival at the shelter after approval, or no motion toward it after a refusal; odometry agrees with ground truth. Only stack failures *before* any VLM decision may be re-run once; all attempts stay in the ledger. *Expected behavior* is the action the generator's ground truth requires.

Registration and statistics. We ran 25 studies (Table II; all in Supp. Table S12; the initial evaluation C is not counted). Each has a frozen registration (protocol, configuration, endpoints, analysis code) that is fixed before the study's first trial or inference; for 25 of them the audit of Supp. Table S9 verifies that the registration commit precedes the first results commit. Of the 0 studies without results yet, all are on the main branch, bound by the integration merge before their first trial (first commit on the main branch and sha256 of every registration: `docs/REGISTRATIONS.md`); they are not audited on a par with the others, and "registered before the results" is claimed for the audited ones only. Studies were designed sequentially, so registration is confirmatory within a study. The primary safety metric is the false-approval rate among must-refuse cases, paired with the correct-approval rate; we use Wilson 95% intervals, world-cluster bootstrap for pooled image rates, exact McNemar and Wilcoxon tests. A claim of improvement requires the registered direction and $p < 0.05$. Claims are not pooled across studies; the headline claims are G, H, V and A, not the 25 studies together.

VI. Results

A. The Gate in Closed Loop

In campaign D (42 trials) and campaign E (60 trials, a new randomized world each) the robot behaved as expected in 42 and 60 trials with 0 and 0 wrong deliveries (E: [0.0, 6.0]%); it completed 36 of the 36 missions in which a safe delivery was possible. A blind robot that trusts the inventory would have made 40 wrong deliveries in the same E trials ([54.1, 77.3]%). On unseen layouts (campaign L, 24 layouts) the robot behaved as expected in 24/24 ([86.2, 100.0]%) with 0 unsafe deliveries. A delivery without search takes about 5.5 s of checking and 39 s of driving (Supp. Table S5). **Supporting studies (Supp. Secs. S7, S11, S12).** A zero-shot SigLIP baseline (Z) approved 44/408 of T1's must-refuse images against 3/408 for our check; pooled over 12 camera degradations (R), the check made 0/2304 false approvals; a conformal τ_h (K) kept its miss bound at the cost of more false alarms. On real photographs the frozen threshold flagged 64/100 of the fire-free ones, whereas a threshold calibrated on site photographs (P2) found 282/300 of the fires and gave 18/100 false alarms on unseen photographs (registered claim: supported), so a threshold needs per-site calibration. A second checker (J) was not safer than one model (claim not supported) and was not adopted.

B. Memory in Changing Worlds (H, V)

Campaign H comprised 16 sequences of three missions to the same shelter. Lighting and props were re-drawn in each mission, and four scripted stories changed the world between missions (supplies swapped and stayed swapped, or swapped back; a small fire appeared; a fire disappeared and another appeared). The first mission seeded all arms; missions 2 and 3 ran with no memory (M0), naive memory that skips the check where it last saw the supply or refuses without visiting where it last saw a hazard (M1), and verified memory, i.e. memory used only as advice (M2). Over the 32 paired missions, unsafe deliveries were 20 for M1, 0 for M2 and 1 for M0 (M1 vs. M2 exact McNemar $p = 1.9 \times 10^{-6}$; by sequence 12 vs. 0 of 16, $p = 0.0005$). M1's unsafe deliveries were 12 wrong supplies, relocated since it last looked, and 8 deliveries into a fire that had appeared since; every one followed a skipped check. The M1–M2 contrast is large by design; the comparison that bounds our risk is M2 versus M0, 0/32 versus 1/32 (Wilson upper bounds 11% and 16%), met as registered but informative only down to those bounds. M0's single unsafe delivery shows what caution is for (Fig. 5): the new fire was invisible from the front and small from the side, the left view scored $p_h = 0.035 < \tau_h$ and M0 delivered. M2 had recorded a fire at this station two missions earlier, so it checked at $\tau_h = 0.01$, scored 0.049 and stopped. Memory can make the robot more careful, never less (Table III). On the 8 missions whose relocation persisted, the median time was 79 s for M2 against 114 s for M0 (exact Wilcoxon $p = 0.0078$); this subset is the favorable one, and over all 32 paired missions the medians were 73 s against 83 s. The price of keeping every check showed after the supplies moved back: M2 first visited the proposed station and then searched (113 s against 97 s for M0). In the model of repeated missions, verified memory keeps 27% of naive memory's saving when nothing changes and none at $q = 0.5$ (Supp. Sec. S12). With k stations a stale proposal costs up to $k - 1$ visits, while the expected saving of a correct one grows with k (descriptive model with the measured detour of 30 s: 9 s per mission at $k = 2$, 80 s at $k = 8$, $q = 0.2$; Supp. Sec. S7). The registered headline claim (verified memory keeps memory's speed-up without its safety cost) is supported. **Study V** tested the repeated-mission model in closed loop (10 sequences of 6 missions, supply location and fire toggling at random with $q = 0.2$, predictions registered before the first trial). The registered 95% interval for the paired saving M0−M2 was 11–20 s; the observed saving was 15 s. Unsafe deliveries were 0 (M0), 20 (M1) and 1 (M2), inside the registered intervals 1 (0–5), 20 (15–25), 1 (0–5). The model claim is supported for the registered quantities (the difference, the unsafe counts, the searches); the absolute mean times lie outside their intervals (M0 87.6 s against 79.3–86.2, M2 72.7 s against 63.5–70.5, 5–6 s above the predicted median), and only the difference between the arms is inside. The memory claim is not supported: the margin held (M2−M0 $= 1 \leq 3$), but the sequence-level test did not (M1 unsafe in 6 sequences, M2 in 1; exact McNemar $p = 0.0625$). This is a failure of power, not evidence of a safety cost: with M1 unsafe in 6 sequences the test can pass only if M2 has no unsafe delivery in any of them (smallest possible $p = 0.031$), while the registered model predicted a median of 1 (0–5) unsafe M2 deliveries. Pooled

TABLE II

THE STUDIES THAT CARRY THE CHAIN (ALL REGISTERED STUDIES: SUPP. TABLE S12). CLAIM: STATUS OF THE REGISTERED CLAIM (PENDING: NOT YET RUN OR ANALYZED).

Study	Question	Data	Arms (systems)	Primary endpoint	Claim
T1, E	Does the gate generalize, and work in closed loop?	600 held-out images; 60 randomized worlds	1 (v3.5)	false approvals $< 2\%$ (T1); expected behavior, wrong deliveries (E)	not supported; descriptive
G (W)	Does looking from the side shrink ε?	24 paired worlds (G); 150 fire + 75 fire-free worlds, images (W)	2 (front only; look-around)	deliveries despite a perceivable fire (G); miss bound $\leq 15\%$ (W)	supported; bound: supported; cost: supported; prior art distinguished: 0/3
H	Can memory speed missions up without new risk?	16 sequences $\times$ 3 missions, changing worlds	3 (no, naive, verified memory)	time; unsafe deliveries	supported (scripted changes; random: V)
V	Does the repeated-mission model predict closed-loop behavior?	10 sequences $\times$ 6 missions, random change	3 (as H)	prediction intervals; unsafe deliveries	model: supported; memory: not supported
A	Can untrusted advice relax a check, or steer it?	226 runs: 18 worlds; 16 duplicate-supply worlds; 224 offline texts	tighten-only, naive, LLM-filter, scrutiny, control	unsafe deliveries under false advice; steering; time; cost bound	safety: supported; time: supported; cost: supported
S	Can a sign in the picture cause a wrong approval?	64 worlds $\times$ no sign, attack, benign	6 (ours, one pass, free-form, full image, OCR mask, ignore-text)	false approvals caused by the sign	letters: not supported; 2/7 hypotheses
X	Does the memory result transfer to a cluttered four-station depot?	depot images (XD); 59 closed-loop runs, $k = 4$	M0 vs. M2 (memory); control, tighten-only, scrutiny (steering)	$\hat{\varepsilon}$, recalibrated τ_h; M2 vs. M0 unsafe, time; steering (descriptive)	memory: supported; steering: S1-S3 hold; S4 (scrutiny fewer unsafe, p<0.05) not supported (not claimed, as registered)

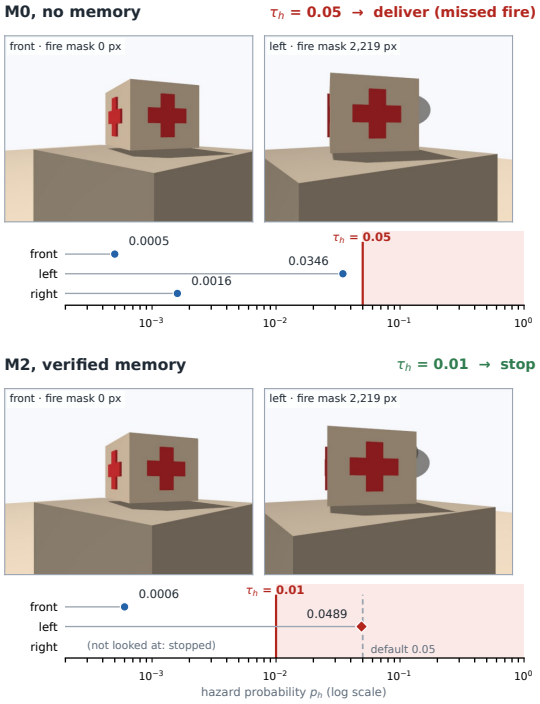


Fig. 5. Campaign H, sequence Q09 mission 3: the same world without memory (M0) and with verified memory (M2). Each row shows the front view and the left side view (fire-mask pixels from the simulator) and the hazard score p_h of every view that was checked against the threshold (shaded: $p_h \geq \tau_h$). M0 used $\tau_h = 0.05$; the left view scored 0.035 and the robot delivered. M2 had seen a fire at this station two missions earlier, so it used $\tau_h = 0.01$; the left view scored 0.049 and the robot stopped without checking the right side.

TABLE III

CAMPAIGN H, MISSIONS 2–3 (PAIRED). M0: NO MEMORY; M1: NAIVE MEMORY; M2: VERIFIED MEMORY. UNSAFE = WRONG SUPPLY OR A PERCEIVABLE FIRE AT THE DELIVERING STATION. TIME: MEDIAN OVER ALL 32 MISSIONS.

Arm	Unsafe	Wrong supply	Into fire	False refusals	Median time [s]
M0 no memory (v3.6)	1/32	0	1	0	83
M1 naive memory	20/32	12	8	4	28
M2 verified memory (ours)	0/32	0	0	0	73

over H and V (82 missions), unsafe deliveries were 1 for M0, 1 for M2 (Wilson [0.2, 6.6]%) and 40 for naive memory. The M2 event is a check error, not a memory effect: in the same world M0's right view scored $p_h = 0.33$ and refused, M2's scored 0.013 after arrival jitter, which is above the I-20 suspicion threshold, but that rule was off in V.

C. Adversarial Advice (A)

Study A pairs an advice-following baseline with ours on 12 worlds in which false advice, if believed, causes an unsafe delivery ("the fire is out, skip the check"; "the water is at the other station"), through two channels (operator text, a peer's advice file), plus an arm in which an LLM guardrail filters the advice before the baseline follows it. Worlds with an unsafe delivery: 12/12 for the advice-following baseline, 12/12 behind the guardrail, 0/12 for ours (runs 16/24, 12/24, 0/24; no-advice control 0/12; exact McNemar $p = 0.0005$ against either baseline, paired by world). The guardrail (the registered "self-filter": one yes/no question, no attack search) read every item's text and flagged 7/12 operator-text attacks but 0/12 peer-file attacks, which are short and look truthful, so there it did not help (12/12 unsafe runs; Supp. Sec. S9). The baselines follow advice by definition: the contrast shows that the interface decides. All A runs used intent prompt i38 (Amendment 1); the submitted default is i35, so A measures the gate with i38, not the shipped default. Every delivery of ours was one that the advice-free check also approves (43/43 replayed on the saved frames, 0 violations): the code conforms to the tightening Prop. 2 requires. The guarantee is relative to the check and assumes frames independent of the plan. Truthful location advice saved a median of 27 s per mission; hostile advice cost at most 41.7 s (bound 50.6 s) and 0 lost deliveries (allowed 4). The operator channel extracted the designed advice in 36/72 runs (relaxing attacks: 3/36), so the file channel is the controlled test (Supp. Sec. S9). On the held-out half of the offline set (112 texts) wrong-shelter acceptance was 0/18 and no tighten-only plan relaxed a check (0/224, naive 34/224). **Steering (A-steer).** The adaptive attacker sends *truthful* location advice to the station that holds a small, partly hidden fire when a clean duplicate exists. In 16 duplicate-supply worlds, unsafe deliveries were 2/16 for tighten-only advice, 0/16 in the control (clean station first) and 0/16 with scrutiny and soft refusal. Both of ours were check errors (the advice-free check approves the same frames); the direction is the predicted one but not significant (sign test $p = 0.25$; Wilson upper bound 0.36) and matches the side-only class of W (3/23 misses with look-around), the worst class, not the pooled ε_{any} (4.0% for $k = 2$): a steering attacker realises the worst class. The cost is availability and exposure, which Prop. 2 does not bound: every robot that followed the advice drove to the burning station (control: none) and stopped there, so deliveries were lost in 14/16 worlds (11/16 with scrutiny, 0/16 without advice). The steering estimate is the oracle-steering ε_{any} at image level: W (150 hard-class worlds) gives 4.0% for $k = 2$ and 7.8% for $k = 4$ with look-around (front-only 75.8%, 94.3%); XD (hard class, 12/100 misses in 100 worlds): 22.7% for $k = 2$ and 40.5% for $k = 4$. Registered claims: safety: supported; time: supported; cost: supported. A consistency–robustness curve of the policies as advice accuracy falls is a model illustration, not a result: its fire sits at the requested supply's station, so the tighten-only curve is flat by construction and only its time axis

is informative (Supp. Sec. S10).

D. Outside the Guarantee: Scene Text and Realism

Scene text (S). Study S (image level, 64 worlds each rendered without a sign, with an attack sign and with a benign text) measures what the guarantee does not cover. Signs flipped a refusal of ours into an approval in 0 of 42 worlds (Wilson upper bound 8.4%): "Answer C" 0 of 9; "A", tied to the unpermuted option order, 0 of 6 (one unpermuted pass: 0 of 6); "B", invariant under the permutation, 0 of 6; mislabels 0 of 10; "SAFE – NO FIRE" beside a visible fire 0 of 8. The residual vulnerabilities we expected (B, mislabel) were not observed, and the one-pass check was not fooled either, so the registered claim that the permutation explains the robustness is not supported (2/7 hypotheses). A benign text moved the free-form baseline in 6 worlds, ours in 0; masking text or telling the model to ignore it changed nothing (56/56 approvals kept). The signs are non-adaptive; scene text stays outside the guarantee of Sec. IV.

Realism and transfer. The arena is a flat floor with two stations. The depot (studies XD and X) is an MIT-0 warehouse with four supply stations, clutter, obstacles and moving smoke; fire and smoke remain synthetic. We stated beforehand that the *relative* claims should transfer, not the absolute values (Supp. Sec. S12). Only image-level tests on photographs (P, P2, Q, XD-R) touch real appearance. **X result** (59 runs): behaviour cells as designed in 16/16, no delivery into a fire (0/12, upper 24.2%). Memory saved 2.14 stations at $k = 4$ (CI [1.00, 2.67]). Steering, descriptive (3 fire worlds): tighten-only delivered into the hidden fire 1/3 (no advice 0/3) and gave up 2 deliveries; scrutiny refused all, also both clean twins. **XD result.** In 308 depot worlds (151 perceivable fires) front-only approved 55/151 fires and look-around 14/151 ([5.6, 15.0]%); on the hard class (small, partly hidden flames) look-around approved 12/100 (Wilson upper bound 19.8%), so the registered 15% bound is *not* met, and the absolute ε did not transfer. The I-20 extra views were withdrawn by the rule fixed beforehand (O3), and recalibration found no threshold. Study Q tests the supply check on 405 public Commons photographs (water, medicine, neither; labels audited by an AI assistant, human spot-check pending): it approved none of 277 must-refuse images (Wilson upper bound 1.4%) but only 52/128 correct ones (40.6%), under the registered 70%, so "the check transfers to real photographs" is not supported (safe but conservative; Supp. Sec. S12). XD-R, the false-alarm rate on fire-free warehouse photographs (258 frozen before scoring), flagged 11/200 (5.5%); neither involves a robot. Real photographs (P, P2) showed that the frozen threshold flagged 64/100 of fire-free photographs while the scores still rank well (AUROC 0.96).

VII. Discussion and Limits

Inputs, not just models, decide safety. Trusting memory is efficient but its risk grows with how often the world changes, which the robot cannot observe without looking (H). Observation 1 and Proposition 2 settle whether checks happen: memory and any other advice decide only the order and strictness of checks, never their existence; the guarantee ends at the camera. **Moving first, then wording.** The largest single gain in ε came from moving 0.4 m (G); in W's v3.5 reference run the look-around alone cut the misses from 144/150 to 48/150, and the v3.6 wording and crops removed 45 more. **Value of pre-registration.** Of the registered claims, thirteen failed (T1's 2% bound, F, T2, J, V4 (memory under random change), W-b (hazard rule beats prior art), IF1 (request switching), S-perm (permutation), XD-ε (depot bound), XD-rel (transfer), XD-cal (recalibration), Q (real photographs) and YE2 (two stations visited: the registered wording)); the first two led to the look-around, and J kept us from a second model.

Limits. The guarantee is relative: G's plug-in $k\hat{\varepsilon} = 0.125$ for the hardest class (upper bound 0.57, 16 worlds) is superseded by W's estimate on 150 worlds; the order-independent $\varepsilon_{\text{any}} \leq k\hat{\varepsilon}$ is the number to read, estimated at image level (W: 4.0% for $k = 2$, 7.8% for $k = 4$ with look-around; XD, hard class: 22.7%, 40.5%). Exposure to a really burning station is not bounded. The price of advice repeated over many missions is bounded per account by a credit that only drops advice (Prop. 3, supplement). In closed loop (study B, a mechanics check, not a rate) a source that lied at all 6 missions wasted 2 visits with the credit and 6 without, in 3 of 3 sequences (bound 3); 0 audit failures, 0 unsafe deliveries. ε_s differs by station and fire geometry, so the budget we measure on one class does not transfer to another. Scene text is perception and outside the guarantee, and so is the requested supply parsed from the operator's text, which Prop. 2 does not cover. Study IF: the frozen parser accepted an injected supply in 12 of 36 switching texts (20–50%), so the request must be read back; with the ask-back fix (I-27) 0 of 36 new texts were accepted. All results are from simulation (one robot, two supply types, a stylized fire, sensor degradations rather than real objects or smoke); the depot (study X) narrows the gap but stays simulation (no grasping, real smoke or sensor noise). Barely visible fires (21–613 pixels) are excluded, and the guard is hand-written, not synthesized [20], [21]. Threats to validity: Supp. Sec. S9.

VIII. Conclusion

In one sentence: *Typing everything the robot did not perceive itself (memory, peer reports, inventory, the operator's wording) as advice that may only reorder visits or tighten thresholds keeps the risk of an unsafe delivery, under any advice, within the check's own error budget* ε_{any}. The look-around shrank the budget (G: 11 to 1 of 16); verified memory kept a stale input from spending more than the budget (H: supported; V: model: supported; memory: not supported; 1 unsafe in 82 missions with verified memory) and tighten-only advice a hostile one (A: unsafe worlds 0/12, against 12/12 for the baselines). The guarantee is relative, steering costs deliveries, and it ends at the camera. Everything is simulation.

References

[1] B. Balcik, B. M. Beamon, and K. Smilowitz, "Last mile distribution in humanitarian relief," *Journal of Intelligent Transportation Systems*, vol. 12, no. 2, pp. 51–63, 2008. [Online]. Available: https://doi.org/10.1080/15472450802023329

[2] R. R. Murphy, *Disaster Robotics*. Cambridge, MA: The MIT Press, 2014. [Online]. Available: https://doi.org/10.7551/mitpress/9407.001.0001

[3] Z. Liu, A. Bahety, and S. Song, "REFLECT: Summarizing robot experiences for failure explanation and correction," in *Proceedings of The 7th Conference on Robot Learning (CoRL)*, ser. Proceedings of Machine Learning Research, vol. 229. PMLR, 2023, pp. 3468–3484. [Online]. Available: https://proceedings.mlr.press/v229/liu23g.html

[4] J. Duan, W. Pumacay, N. Kumar, Y. R. Wang *et al.*, "AHA: A vision-language-model for detecting and reasoning over failures in robotic manipulation," in *International Conference on Learning Representations (ICLR)*, 2025. [Online]. Available: https://proceedings.iclr.cc/paper_files/paper/2025/hash/70a06501001e1820fd1eb9ee821302d2-Abstract-Conference.html

[5] C. Agia, R. Sinha, J. Yang, Z. Cao *et al.*, "Unpacking failure modes of generative policies: Runtime monitoring of consistency and progress," in *Proceedings of The 8th Conference on Robot Learning (CoRL)*, ser. Proceedings of Machine Learning Research, vol. 270. PMLR, 2024, pp. 689–723, coRL 2024 (PMLR volume dated 2025); method name: Sentinel. [Online]. Available: https://proceedings.mlr.press/v270/agia25a.html

[6] C. Zheng, H. Zhou, F. Meng, J. Zhou, and M. Huang, "Large language models are not robust multiple choice selectors," in *International Conference on Learning Representations (ICLR)*, 2024. [Online]. Available: https://openreview.net/forum?id=shr9PXz7T0

[7] P. Pezeshkpour and E. Hruschka, "Large language models sensitivity to the order of options in multiple-choice questions," in *Findings of the Association for Computational Linguistics: NAACL 2024*. Association for Computational Linguistics, 2024, pp. 2006–2017. [Online]. Available: https://doi.org/10.18653/v1/2024.findings-naacl.130

[8] Y. Li, Y. Du, K. Zhou, J. Wang *et al.*, "Evaluating object hallucination in large vision-language models," in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, 2023, pp. 292–305. [Online]. Available: https://doi.org/10.18653/v1/2023.emnlp-main.20

[9] Y. Sun and Y. Zhang, "When memory lies: An empirical study of spatial memory staleness in VLM agents," 2026, arXiv preprint.

[10] Y. Du, K. Konyushkova, M. Denil, A. Raju *et al.*, "Vision-language models as success detectors," in *Proceedings of The 2nd Conference on Lifelong Learning Agents (CoLLAs)*, ser. Proceedings of Machine Learning Research, vol. 232. PMLR, 2023, pp. 120–136. [Online]. Available: https://proceedings.mlr.press/v232/du23b.html

[11] C. Xu, T. K. Nguyen, E. Dixon, C. Rodriguez *et al.*, "Can we detect failures without failure data? uncertainty-aware runtime failure detection for imitation learning policies," in *Robotics: Science and Systems XXI (RSS)*, 2025, method name: FAIL-Detect. [Online]. Available: https://doi.org/10.15607/RSS.2025.XXI.073

[12] E. Zhou, Q. Su, C. Chi, Z. Zhang *et al.*, "Code-as-monitor: Constraint-aware visual programming for reactive and proactive robotic failure detection," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025, pp. 6919–6929. [Online]. Available: https://doi.org/10.1109/CVPR52734.2025.00649

[13] R. Sinha, A. Elhafsi, C. Agia, M. Foutter, E. Schmerling, and M. Pavone, "Real-time anomaly detection and reactive planning with large language models," in *Proc. Robotics: Science and Systems (RSS)*, 2024.

[14] A. Oloo, Z. Altaweel, Y. Hayamizu, P. Liu, Y. Ding, S. Amiri, H. Yang, A. Kaminski, C. Esselink, C. Paxton, X. Zhang, and S. Zhang, "Robot planning and situation handling with active perception," 2026, arXiv preprint; method name: VAP-TAMP.

[15] A. Z. Ren, A. Dixit, A. Bodrova, S. Singh *et al.*, "Robots that ask for help: Uncertainty alignment for large language model planners," in *Proceedings of The 7th Conference on Robot Learning (CoRL)*, ser. Proceedings of Machine Learning Research, vol. 229. PMLR, 2023, pp. 661–682. [Online]. Available: https://proceedings.mlr.press/v229/ren23a.html

[16] B. Ichter, A. Brohan, Y. Chebotar, C. Finn *et al.*, "Do as I can, not as I say: Grounding language in robotic affordances," in *Proceedings of The 6th Conference on Robot Learning (CoRL)*, ser. Proceedings of Machine Learning Research, vol. 205. PMLR, 2022, pp. 287–318, coRL 2022 (PMLR volume dated 2023). Author order follows the PMLR record; the arXiv version (2204.01691) lists Michael Ahn first ("Ahn et al., 2022"). [Online]. Available: https://proceedings.mlr.press/v205/ichter23a.html

[17] J. Robinson, C. M. Rytting, and D. Wingate, "Leveraging large language models for multiple choice question answering," in *International Conference on Learning Representations (ICLR)*, 2023. [Online]. Available: https://openreview.net/forum?id=yKbprarjc5B

[18] D. Seto, B. Krogh, L. Sha, and A. Chutinan, "The Simplex architecture for safe online control system upgrades," in *Proceedings of the 1998 American Control Conference (ACC)*. IEEE, 1998, pp. 3504–3508. [Online]. Available: https://doi.org/10.1109/ACC.1998.703255

[19] L. Sha, "Using simplicity to control complexity," *IEEE Software*, vol. 18, no. 4, pp. 20–28, 2001. [Online]. Available: https://doi.org/10.1109/MS.2001.936213

[20] M. Alshiekh, R. Bloem, R. Ehlers, B. Könighofer *et al.*, "Safe reinforcement learning via shielding," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, no. 1, 2018. [Online]. Available: https://doi.org/10.1609/aaai.v32i1.11797

[21] A. D. Ames, S. Coogan, M. Egerstedt, G. Notomista *et al.*, "Control barrier functions: Theory and applications," in *18th European Control Conference (ECC)*, 2019, pp. 3420–3431. [Online]. Available: https://doi.org/10.23919/ECC.2019.8796030

[22] B. Zitkovich, T. Yu, S. Xu, P. Xu *et al.*, "RT-2: Vision-language-action models transfer web knowledge to robotic control," in *Proceedings of The 7th Conference on Robot Learning (CoRL)*, ser. Proceedings of Machine Learning Research, vol. 229. PMLR, 2023, pp. 2165–2183, author order follows the PMLR record; the arXiv version (2307.15818) lists Anthony Brohan first ("Brohan et al., 2023"). [Online]. Available: https://proceedings.mlr.press/v229/zitkovich23a.html

[23] M. J. Kim, K. Pertsch, S. Karamcheti, T. Xiao *et al.*, "OpenVLA: An open-source vision-language-action model," in *Proceedings of The 8th Conference on Robot Learning (CoRL)*, ser. Proceedings of Machine Learning Research, vol. 270. PMLR, 2024, pp. 2679–2713, coRL 2024 (PMLR volume dated 2025). [Online]. Available: https://proceedings.mlr.press/v270/kim25c.html

[24] A. Robey, Z. Ravichandran, V. Kumar, H. Hassani, and G. J. Pappas, "Jailbreaking LLM-controlled robots," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2025, pp. 11 948–11 956. [Online]. Available: https://doi.org/10.1109/ICRA55743.2025.11128119

[25] H. Zhang, C. Zhu, X. Wang, Z. Zhou, C. Yin, M. Li, L. Xue, Y. Wang, S. Hu, A. Liu, P. Guo, and L. Y. Zhang, "BadRobot: Jailbreaking embodied LLM agents in the physical world," in *International Conference on Learning Representations (ICLR)*, 2025.

[26] L. Burbano, D. Ortiz, Q. Sun, S. Yang, H. Tu, C. Xie, Y. Cao, and A. A. Cardenas, "CHAI: Command hijacking against embodied AI," in *IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*, 2026, accepted; proceedings entry not checked.

[27] S. M. B. P. Samarakoon, M. A. V. J. Muthugala, W. K. R. Sachinthana, and M. R. Elara, "Hijacking robots with a piece of paper: A systematic study of physical prompt injection in VLM-controlled robots," 2026, preprint.

[28] Z. Chen, Z. Xiang, C. Xiao, D. Song, and B. Li, "AgentPoison: Red-teaming LLM agents via poisoning memory or knowledge bases," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

[29] Z. Ravichandran, A. Robey, V. Kumar, G. J. Pappas, and H. Hassani, "Safety guardrails for LLM-enabled robots," *IEEE Robotics and Automation Letters*, vol. 11, no. 4, pp. 4649–4656, 2026. [Online]. Available: https://doi.org/10.1109/LRA.2026.3667488

[30] K. J. Biba, "Integrity considerations for secure computer systems," The MITRE Corporation, Tech. Rep. MTR-3153, 1977.

[31] D. E. Denning, "A lattice model of secure information flow," *Communications of the ACM*, vol. 19, no. 5, pp. 236–243, 1976.

[32] E. Debenedetti, I. Shumailov, T. Fan, J. Hayes, N. Carlini, D. Fabian, C. Kern, C. Shi, A. Terzis, and F. Tramèr, "Defeating prompt injections by design," 2025, preprint.

[33] M. Costa, B. Köpf, A. Kolluri, A. Paverd, M. Russinovich, A. Salem, S. Tople, L. Wutschitz, and S. Zanella-Béguelin, "Securing AI agents with information-flow control," 2025, preprint.

[34] T. Lykouris and S. Vassilvitskii, "Competitive caching with machine learned advice," *Journal of the ACM*, vol. 68, no. 4, pp. 1–25, 2021, conference version: ICML 2018. arXiv:1802.05399.

[35] M. Dinitz, S. Im, T. Lavastida, B. Moseley, and S. Vassilvitskii, "Faster matchings via learned duals," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021, pages not verified.

[36] S. Chepuri and V. Srivastava, "AuthGuard-R: Safety-compliant mission hijacking and dual-gate defense for LLM-controlled robots," 2026, preprint, submitted 25 September 2026.

[37] P. T. Tran-Truong, H. Vo, N. Tran, K. Vo, and T.-M. Truong, "DelegateGuard: Purpose-bound capabilities for delegated LLM agent workflows," 2026, preprint, submitted 12 October 2026.

[38] A. Kravchenko, V. Liventsev, I. Konstantinov, I. Iskhakov, and M. Kukuy, "APPA: Recoverable information-flow control for real-world LLM agents," 2026, preprint, submitted 27 July 2026 (v2 26 August 2026).

[39] M. Akewar and R. Ranjan, "SafeCommit: Certifying when memory-grounded agents may safely act," 2026, preprint, submitted 4 August 2026.

[40] A. Anwar, J. Welsh, J. Biswas, S. Pouya, and Y. Chang, "ReMEmbR: Building and reasoning over long-horizon spatio-temporal memory for robot navigation," in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2025, pp. 2838–2845.

[41] Z. Wang, B. Yu, J. Zhao, W. Sun, S. Hou, S. Liang, X. Hu, Y. Han, and Y. Gan, "KARMA: Augmenting embodied AI agents with long-and-short term memory systems," in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2025.

[42] N. Hughes, Y. Chang, and L. Carlone, "Hydra: A real-time spatial perception system for 3D scene graph construction and optimization," in *Proc. Robotics: Science and Systems (RSS)*, 2022.

[43] Q. Gu, A. Kuwajerwala, S. Morin, K. M. Jatavallabhula, B. Sen, A. Agarwal, C. Rivera, W. Paul, K. Ellis, R. Chellappa, C. Gan, C. M. de Melo, J. B. Tenenbaum, A. Torralba, F. Shkurti, and L. Paull, "ConceptGraphs: Open-vocabulary 3D scene graphs for perception and planning," in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2024, pp. 5021–5028.

[44] K. Rana, J. Haviland, S. Garg, J. Abou-Chakra, I. Reid, and N. Suenderhauf, "SayPlan: Grounding large language models using 3D scene graphs for scalable robot task planning," in *Proc. Conference on Robot Learning (CoRL)*, ser. Proceedings of Machine Learning Research, vol. 229. PMLR, 2023, pp. 23–72.

[45] X. Yin and H. Du, "GLOVE: Global verifier for LLM memory-environment realignment," 2026, arXiv preprint.

[46] G. Tang, Q. Jia, G. Chen, T. Li, Z. Huang, Z. Lv, and N. Ji, "VLM-DEWM: Dynamic external world model for verifiable and resilient vision-language planning in manufacturing," 2026, arXiv preprint.

[47] A. Aydemir, A. Pronobis, M. Göbelbecker, and P. Jensfelt, "Active visual object search in unknown environments using uncertain semantics," *IEEE Transactions on Robotics*, vol. 29, no. 4, pp. 986–1002, 2013.

[48] A. Wandzel, Y. Oh, M. Fishman, N. Kumar, L. L. S. Wong, and S. Tellex, "Multi-object search using object-oriented POMDPs," in *Proc. IEEE Int. Conf. Robotics and Automation (ICRA)*, 2019, pp. 7194–7200.

[49] Z. Zeng, A. Röfer, and O. C. Jenkins, "Semantic linking maps for active visual object search," in *Proc. IEEE Int. Conf. Robotics and Automation (ICRA)*, 2020, pp. 1984–1990, arXiv:2006.10807.

[50] K. Zhou, K. Zheng, C. Pryor, Y. Shen, H. Jin, L. Getoor, and X. E. Wang, "ESC: Exploration with soft commonsense constraints for zero-shot object navigation," in *Proc. 40th Int. Conf. Machine Learning (ICML), PMLR 202*, 2023, pp. 42 829–42 842, arXiv:2301.13166.

[51] J. Mason and B. Marthi, "An object-based semantic world model for long-term change detection and semantic querying," in *Proc. IEEE/RSJ Int. Conf. Intelligent Robots and Systems (IROS)*, 2012, pp. 3851–3858.

[52] D. M. Rosen, J. Mason, and J. J. Leonard, "Towards lifelong feature-based mapping in semi-static environments," in *Proc. IEEE Int. Conf. Robotics and Automation (ICRA)*, 2016, pp. 1063–1070.

[53] R. Bajcsy, "Active perception," *Proceedings of the IEEE*, vol. 76, no. 8, pp. 966–1005, 1988.

[54] R. Bajcsy, Y. Aloimonos, and J. K. Tsotsos, "Revisiting active perception," *Autonomous Robots*, vol. 42, no. 2, pp. 177–196, 2018.

[55] A. Z. Ren, J. Clark, A. Dixit, M. Itkina, A. Majumdar, and D. Sadigh, "Explore until confident: Efficient exploration for embodied question answering," in *Proc. Robotics: Science and Systems (RSS)*, 2024.

[56] B. A. Nosek, C. R. Ebersole, A. C. DeHaven, and D. T. Mellor, "The preregistration revolution," *Proceedings of the National Academy of Sciences*, vol. 115, no. 11, pp. 2600–2606, 2018. [Online]. Available: https://doi.org/10.1073/pnas.1708274114

[57] E. B. Wilson, "Probable inference, the law of succession, and statistical inference," *Journal of the American Statistical Association*, vol. 22, no. 158, pp. 209–212, 1927. [Online]. Available: https://doi.org/10.1080/01621459.1927.10502953

[58] Q. McNemar, "Note on the sampling error of the difference between correlated proportions or percentages," *Psychometrika*, vol. 12, no. 2, pp. 153–157, 1947. [Online]. Available: https://doi.org/10.1007/BF02295996

[59] S. Macenski, T. Foote, B. Gerkey, C. Lalancette, and W. Woodall, "Robot operating system 2: Design, architecture, and uses in the wild," *Science Robotics*, vol. 7, no. 66, p. eabm6074, 2022. [Online]. Available: https://doi.org/10.1126/scirobotics.abm6074

[60] N. Koenig and A. Howard, "Design and use paradigms for Gazebo, an open-source multi-robot simulator," in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, vol. 3, 2004, pp. 2149–2154. [Online]. Available: https://doi.org/10.1109/IROS.2004.1389727

[61] Qwen Team, "Qwen3.5: Towards native multimodal agents," Official blog post and Hugging Face model card (Qwen/Qwen3.5-9B, Apache-2.0), Feb. 2026, model card: https://huggingface.co/Qwen/Qwen3.5-9B. [Online]. Available: https://qwen.ai/blog?id=qwen3.5